


lazy vs rich, $N = 256$

feature-kernel drift



training loss

