


# lazy vs rich, $N = 256$

## feature-kernel drift



## training loss

